
Liar, Liar: Beyond Vocabulary Suppression

Aryan Gupta¹

¹University of Illinois Urbana-Champaign

aryang9@illinois.edu

github.com/aryan-cs/liar-liar

Activation-steering vectors for honesty raise truthfulness benchmarks, and this is widely read as evidence that the intervention moves an upstream concept. We ask whether the effect instead rides the model’s output readout: a bias toward honest-coded words rather than a change in what the model computes. We separate the two by projecting a steering vector onto the orthogonal complement of the effective-unembedding rows (the unembedding composed with the RMSNorm Jacobian) for a chosen set of honesty-coded tokens. The projected vector has certified zero direct logit effect on those tokens, so any behavior it preserves must travel the indirect path through downstream layers; the preserved fraction ρ measures depth. The statistic is only meaningful once two pitfalls are controlled. A steering magnitude must be selected under a coherence gate, because an over-large coefficient inflates teacher-forced multiple-choice scores while destroying free generation and turning ρ into a ratio of noise; and the unembedding subspace must beat a random subspace of equal rank, or its removal proves nothing. On Llama-3-8B-Instruct we decompose two honesty-steering constructions head to head and find a double dissociation. The popular contrastive (CAA) vector shifts honesty-coded vocabulary by over a logit without improving truthfulness at any coherent strength; its apparent benchmark gain arises only after generation has collapsed (held-out perplexity $68.97\times$ baseline), where the naive analysis also reports a confident and meaningless $\rho \approx 1$. The mass-mean vector improves truthfulness at full coherence, with a test-set MC2 gain of $+0.073$, and its effect survives certified excision of its entire direct vocabulary readout: $\rho = 0.95 [+0.79, +1.13]$, indistinguishable from random-subspace controls, never below 0.94 at any excision rank from $k=16$ to $k=1024$, and intact under paraphrase ($\rho_{\text{OOD}} = 0.97$), while the logit lens shows downstream layers rebuilding the suppressed readout within two blocks. Where honesty steering works it is not vocabulary suppression, and where it is vocabulary-level it does not work. Code, token sets, certificates, and a formal proof of the construction are released.

1 Introduction

Activation steering has become a standard lever for controlling language-model behavior: add a contrastively constructed vector to the residual stream at a middle layer, and benchmark measures of honesty, refusal, or sentiment shift in the intended direction (Zou et al., 2023; Rinsky et al., 2024; Li et al., 2023; Arditi et al., 2024). The appeal is mechanistic. The intervention appears to reach in and adjust a high-level concept the model uses, and for honesty in particular this reading underwrites a safety story: if a steering vector moves an internal representation of truthfulness, the same vector should generalize across phrasings, domains, and languages.

A deflationary alternative is rarely ruled out. An honesty vector that merely tilts the output distribution toward words like *true* and *honest* and away from *lie* and *deceive* would post the same

benchmark gain while changing nothing upstream. The two accounts, an upstream concept versus a readout-level vocabulary bias, predict the same TruthfulQA score and very different out-of-distribution behavior. The standard evaluation does not separate them, because adding a vector at layer ℓ^* moves both the residual the unembedding reads directly and the residual that downstream layers transform first.

We separate the two paths by construction. The first-order logit response to an injected vector v splits into a direct term $\widetilde{W}_U^* v$, the part the readout reads unchanged, and an indirect term $\widetilde{W}_U^* M v$, the part that travels through downstream attention and MLP blocks (Section 3). For a set T of honesty-coded tokens, we project the steering vector onto the orthogonal complement of the corresponding rows of the *effective* unembedding, meaning the unembedding composed with the RMSNorm Jacobian at the readout point. The projected vector $v^\perp$ has exactly zero direct effect on every token in T , certified to machine precision, so any behavior it preserves must travel the indirect path. The fraction preserved, $\rho = \Delta(v^\perp)/\Delta(v_{\text{dec}})$, is a depth statistic: near 0 if the effect was direct readout, near 1 if it was downstream computation.

Two design choices make the statistic trustworthy, and both turned out to matter. First, the ratio is only meaningful when its denominator, the steering effect itself, is real. We select each vector’s operating point under a coherence gate that rejects magnitudes which inflate held-out perplexity, because a large steering coefficient can raise teacher-forced multiple-choice scores while reducing free generation to noise; a sweep that ignores this picks a setting where ρ is a ratio of two noise terms. Second, the unembedding subspace must be shown to be special: we compare projecting out the aligned T against projecting out a random subspace of equal rank. If a random projection erases the effect just as well, the effect was never concentrated in the readout, and a high ρ means nothing.

Contributions. (i) A token-conditional projection that zeroes a steering vector’s direct logit contribution on a chosen token set, with the RMSNorm correction that prior unembedding-orthogonal constructions omit, and machine-precision zero-direct-effect certificates. (ii) A depth statistic ρ with a coherence-gated operating-point protocol and a random-projection control that together make it interpretable, on Llama-3-8B-Instruct. (iii) A head-to-head decomposition of two honesty-steering constructions showing a double dissociation: contrastive activation addition moves honesty-coded vocabulary without moving truthfulness at any coherent operating point, while the mass-mean direction’s real gain (a test-set ΔMC2 of $+0.073$ [$+0.038$, $+0.108$]) survives certified excision of its direct vocabulary readout ($\rho = 0.95$ [$+0.79$, $+1.13$]), indistinguishably from random-subspace controls, in and out of distribution. All code, token lists, certificates, and the formal proof are released.

2 Related work

Activation steering. Representation engineering adds contrastively constructed vectors to the residual stream to modulate high-level behavior (Zou et al., 2023); contrastive activation addition (Rimsky et al., 2024) and inference-time intervention (Li et al., 2023) are the canonical honesty-adjacent instantiations, and Arditi et al. (2024) show refusal is mediated by a single direction that can be removed from every writing matrix. Tan et al. (2024) document that steering vectors often generalize poorly out of distribution; our depth statistic offers a mechanism-level account of when that should happen.

Concept erasure. LEACE (Belrose et al., 2023) gives the minimum-norm affine map that removes linear decodability of a concept, following iterative nullspace projection (Ravfogel et al., 2020) and rank-constrained adversarial erasure (Ravfogel et al., 2022). Our projection is structurally the LEACE construction with two changes: the erased subspace is read directly off the model’s effective unembedding rather than learned from labels, and the constraint is zero direct logit contribution on a token set rather than zero probe accuracy.

Unembedding geometry. Park et al. (2024) formalize the duality between context-side and unembedding-side concept directions under a causal inner product; our rank-one variant projects against the unembedding-side honesty direction, and our Mahalanobis-weighted variant recovers their inner product. Marks and Tegmark (2024) construct truth directions by mass-mean differences, which we use both as an alternative steering vector and as the readout-side direction of the rank-one test. The direct logit-attribution path through the residual stream is the lens of Elhage et al. (2021) and nostalgebraist (2020).

Closest precedents. Three works overlap most directly. Venkatesh and Kurupath (2026) prove steering vectors are non-identifiable up to perturbations in the null space of the full activation-to-logit Jacobian; our subspace is the readout-restricted version, and where they establish ambiguity we measure what the ambiguity carries. Nadaf (2026) show function vectors steer behavior the logit lens cannot decode, establishing the off-readout channel in the in-context-learning setting. The unembedding-steering benchmark of van Deventer (2024) implements W_U -orthogonal steering for sentiment on Gemma-2-9B. None applies the construction to honesty steering, none corrects for the RMSNorm Jacobian, and none reports a quantitative depth decomposition with certificates; those are the contributions here.

3 Method

3.1 Setup and notation

Let $\mathcal{M}$ be a decoder-only transformer with L layers, residual width d , and vocabulary size V . The residual stream $h_i^{(\ell)} \in \mathbb{R}^d$ is updated additively by each block, and the logits at the final position are

$$\ell(h_n^{(L)}) = W_U \cdot \text{RMSNorm}_\gamma(h_n^{(L)}), \quad \text{RMSNorm}_\gamma(z) = \gamma \odot \frac{z}{\sqrt{d^{-1}\|z\|_2^2 + \varepsilon}}, \quad (1)$$

with $W_U \in \mathbb{R}^{V \times d}$ the unembedding and $\gamma \in \mathbb{R}^d$ a learned gain (Zhang and Sennrich, 2019). A steering intervention adds a vector αv_{dec} to the residual stream at a middle layer ℓ^* at every position. We use the contrastive mean-difference construction (Rimsky et al., 2024; Zou et al., 2023): given prompts rendered under an honest and a deceptive system prompt, v_{dec} is the difference of mean final-token residuals at ℓ^* .

Effective unembedding. For a perturbation δz at the readout point z , the first-order logit response is governed not by W_U but by the composition with the RMSNorm Jacobian,

$$\widetilde{W}_U(z) := W_U J_\gamma(z), \quad J_\gamma(z) = \frac{1}{\sigma(z)} \text{diag}(\gamma) \left(I_d - \frac{zz^\top}{d\sigma(z)^2} \right), \quad (2)$$

where $\sigma(z) = \sqrt{d^{-1}\|z\|_2^2 + \varepsilon}$. Projecting against raw W_U rows, as in prior unembedding-orthogonal constructions, leaves a residual direct effect through the gain and the rank-one term. We average J_γ over a calibration set of readout points (Section 4) and write $\widetilde{W}_U^*$ for the resulting matrix. All zero-direct-effect certificates are stated with respect to this averaged matrix; their transfer to an individual readout point is governed by the per-point deviation of the Jacobian from its average. That deviation is measured, not assumed: evaluated at each of the 256 calibration points individually, the residual per-point direct effect of the projected vector on its aligned-64 set is at most 0.011 logits per unit α for the mass-mean vector and 0.019 for CAA, under 3% of the unprojected vectors’ median per-point direct effects (0.60 and 0.77 respectively).

3.2 Token-conditional orthogonalization

A global version of the test, a perturbation invisible to *every* token, is impossible: for every model we consider, $V > d$ and W_U has full column rank, so $\ker(W_U) = \{0\}$ (Proposition 3.1 of the

supplementary proof). The construction must restrict the readout to a chosen token set.

Fix a deception-coded token set $T \subset \{1, \dots, V\}$ with $k = |T| \ll d$, and let $A = \widetilde{W}_U^*[T, :] \in \mathbb{R}^{k \times d}$ be the corresponding rows. With A^+ the Moore–Penrose pseudoinverse,

$$P_T^\perp = I_d - A^+ A, \quad v^\perp = P_T^\perp v_{\text{dec}}, \quad v^\parallel = v_{\text{dec}} - v^\perp. \quad (3)$$

By construction $A v^\perp = 0$: the projected vector has exactly zero first-order logit contribution at every token in T . Among all vectors satisfying that constraint, $v^\perp$ is the one closest to v_{dec} in Euclidean norm (the LEACE-style minimum-norm characterization; Theorem 8.1 of the supplementary proof, after Belrose et al., 2023).

3.3 Direct and indirect paths

Writing $M^{(\ell^* \rightarrow L)}$ for the summed Jacobian of all attention and MLP contributions downstream of ℓ^* , the first-order logit response to an injected v decomposes as

$$\Delta \ell = \underbrace{\widetilde{W}_U^* v}_{\text{direct}} + \underbrace{\widetilde{W}_U^* M^{(\ell^* \rightarrow L)} v}_{\text{indirect}} + O(\|v\|^2). \quad (4)$$

For $v = v^\perp$ the direct term vanishes on T identically, so any change in the T -restricted logits (and, to the extent that behavior on a truthfulness benchmark is mediated by T , any change in measured behavior) must travel through downstream computation (Theorem 6.1 of the supplementary proof).

3.4 The depth statistic

Let $\beta(v)$ be a benchmark score under intervention v and $\Delta(v) = \beta(v) - \beta(0)$. The depth statistic is

$$\rho(v_{\text{dec}}, T) = \frac{\Delta(v^\perp)}{\Delta(v_{\text{dec}})}. \quad (5)$$

Under a pure vocabulary-suppression account, $\rho \approx 0$ once T exhausts the readout-aligned mass of v_{dec} ; under a pure upstream-concept account, $\rho \approx 1$. We report ρ with paired bootstrap confidence intervals, its stability σ_T across independent constructions of T , and a per-prompt analogue ρ_η defined on the honest-shift

$$\eta(v) = [\mu_{T^+}(v) - \mu_{T^-}(v)] - [\mu_{T^+}(0) - \mu_{T^-}(0)], \quad (6)$$

where $\mu_{T^\pm}$ is the mean logit over honest-coded (T^+) and deceptive-coded (T^-) tokens at the answer position.

Token-set constructions. We instantiate T three ways. *Curated*: a fixed lexicon of 32 honest and 32 deceptive lemmas expanded to single-token surface variants. *Statistical*: the top-32 tokens per side by mean first-position logit shift between honest- and deceptive-prompted contexts on held-out prompts. *Aligned- k* : the top- k tokens by $|\widetilde{W}_U^* v_{\text{dec}}|$, the k tokens the steering vector moves most through the direct path. The aligned construction is the strongest form of the test: it removes exactly the direct-readout content of v_{dec} itself, and sweeping k traces how the surviving effect $\rho(k)$ decays as more of the direct path is excised.

Controls. Three controls separate the unembedding subspace from generic perturbation. (i) *Random projection*: project out a random k -dimensional subspace (three seeds), matching the dimension count of the aligned-64 set. (ii) *Norm matching*: rescale $v^\perp$ to $\|v_{\text{dec}}\|$, separating direction from magnitude. (iii) *Parallel injection*: inject $v^\parallel$ alone, the readout-aligned rank- $\leq k$ component; under the suppression account this component should carry the behavioral effect.

4 Experimental setup

Model. All experiments use Llama-3-8B-Instruct (Grattafiori et al., 2024) ($L = 32$, $d = 4096$, $V = 128,256$, RMSNorm) in bfloat16. Projections and certificates are computed in float64.

Data and splits. TruthfulQA (Lin et al., 2022) provides the behavioral benchmark in its multiple-choice form (817 questions). A seeded shuffle assigns 120 questions to a validation slice used only for operating-point selection, 200 to construct mass-mean statements, and the remaining 497 to the held-out test slice on which every headline number is computed. The CAA vector is built from 256 Alpaca instructions (Taori et al., 2023) rendered under an honest and a deceptive system prompt; a disjoint block of 96 instructions serves as held-out fluency text for the coherence gate, and the remaining 48 instructions of the 400-instruction pool feed the statistical token set. The statistical sets are data-driven rather than semantic: they admit punctuation and fragment tokens whose first-position logits shift under the system-prompt contrast, and they function as a control construction, not a curated lexicon. The same 256 instructions, re-rendered as plain chat prompts with no system prompt, also serve as the calibration set for the effective unembedding: J_γ is averaged over their pre-norm final-layer residuals at the last prompt position. Vector-construction data and evaluation data share no prompts.

Steering-vector families. We decompose two constructions head to head. The CAA vector (Rimsky et al., 2024) is the difference of mean final-token residuals between honest- and deceptive-prompted renderings, computed per layer. The *mass-mean* vector (Marks and Tegmark, 2024) is the difference of mean residuals between true and false Q:/A: statements assembled from the 200 reserved TruthfulQA questions, computed at the same candidate layers. Each family is analyzed at its own operating point with its own projections; the unembedding apparatus is shared.

Coherence-gated operating points. For each family we sweep $\ell \in \{12, 14\}$ and $\alpha \in \{0.5, 1, 1.5, 2, 3, 4\}$ and score each setting two ways: the validation MC2 improvement, and the ratio of held-out per-token perplexity under steering to baseline on the 96 reserved Alpaca instructions. A setting is *coherent* if its perplexity ratio is at most 1.5; the operating point is the coherent setting with the largest validation MC2 gain. This gate exists because teacher-forced multiple-choice scores are nearly blind to generation collapse: a sufficiently large α raises MC2 while reducing free generation to gibberish (Section 5), and a sweep that maximizes MC2 alone selects exactly that regime. The gate selects ($\ell^*=12$, $\alpha^*=3$) for CAA and ($\ell^*=12$, $\alpha^*=4$) for mass-mean; Figure 1 shows the full grid.

Scoring. TruthfulQA MC follows the original scoring rule: the total log-probability of each answer choice, with MC1 the accuracy of the argmax choice and MC2 the normalized probability mass on true choices. Choices are scored zero-shot under the prompt Q: {question}\nA: rather than behind the original six-example QA primer; every condition shares this prompt, so all reported contrasts are internal, and absolute scores are not directly comparable to primer-based numbers in the literature. Prompt and continuation are tokenized separately and concatenated at the id level; we verified that this equals tokenizing the joined string for all 5,882 prompt-choice pairs in the benchmark (0 mismatches). The intervention is applied at every token position. At the answer position we additionally record the logits of every token in the union of the curated, spillover, statistical, and aligned-64 token sets, from which the honest-shift η and the spillover readouts are computed without further forward passes.

Conditions. For each family, the headline matrix evaluates on the 497 test questions: a shared unsteered baseline; the family’s vector v_{dec} ; $v^\perp$ for aligned- k with $k \in \{16, 64, 256, 1024\}$; $v^\perp$ for the curated and statistical sets; the parallel component $v^\parallel$ (aligned-64); the norm-matched $v^\perp$ (aligned-64); and random 64-dimensional projections (three seeds). All conditions within a family share that family’s (ℓ^*, α^*) .




Figure 1: Coherence calibration: held-out perplexity ratio against steering strength α for each candidate layer, with the gate at 1.5 (dashed). CAA crosses the gate between $\alpha=3$ and $\alpha=4$ at layer 12 and earlier at layer 14; the mass-mean vector stays coherent across nearly the whole grid.

Out-of-distribution probe. The first 150 test questions are paraphrased by the unsteered model (greedy decoding, meaning-preserving instruction) and rescored under baseline, v_{dec} , and $v^\perp$ (aligned-64) per family, giving an OOD depth ratio ρ_{OOD} . The spillover readout evaluates η on stem-disjoint synonym sets that were never projected out. A logit-lens trajectory (nostalggebraist, 2020) tracks the curated- T readout across layers under v_{dec} , $v^\perp$, and $v^\parallel$ to localize where the honest-shift re-emerges.

Reproducibility. The pipeline is staged and checkpointed with fixed seeds; all code, token lists, certificates, and the formal proof are in the repository. The zero-direct-effect certificates $\max_t |(Av^\perp)_t|$ are at most 7×10^{-15} logits across every family and token set (machine precision in float64), against pre-projection direct effects of up to 0.77 logits per unit α . Forward passes run in bfloat16 without enforced kernel determinism, so individual per-question scores shift at the third decimal across re-runs; every reported interval is orders of magnitude wider.

5 Results

5.1 The naive protocol selects a broken instrument

The standard recipe for choosing a steering strength is to sweep layer and coefficient and keep the setting with the largest validation benchmark gain. On our sweep that recipe selects $\ell=10$, $\alpha=8$. At this setting the injected vector is $2.77\times$ the median residual norm at the injection layer, held-out perplexity is $68.97\times$ baseline, and free generation degenerates into repetitive token fragments (Appendix A). The benchmark does not register the collapse. Teacher-forced MC2 nominally improves by $+0.014$ $[-0.040, +0.065]$ while MC1 falls by -0.080 $[-0.133, -0.028]$ from a baseline of 0.312. Under teacher forcing, every answer token is conditioned on the reference text rather than on the model’s own continuation, so a perturbation that destroys generation can still reweight fixed continuations favorably.

condition	MC2	Δ MC2	ρ
CAA ($\ell^*=12$, $\alpha^*=3$, PPL ratio 1.30; baseline MC2 = 0.471)			
v_{dec}	0.436	−0.035	—
$v^\perp$ al-16	0.426	−0.045	1.30 [−0.12, 3.16]
$v^\perp$ al-64	0.429	−0.042	1.20 [0.01, 2.99]
$v^\perp$ al-256	0.424	−0.047	1.35 [−1.10, 4.46]
$v^\perp$ al-1024	0.424	−0.047	1.35 [−1.70, 5.99]
$v^\perp$ cur	0.437	−0.034	0.97 [0.63, 1.25]
$v^\perp$ stat	0.436	−0.035	1.01 [0.75, 1.37]
$v^\parallel$ al-64	0.481	+0.010	−0.30 [−2.63, 0.67]
$v^\perp$ al-64 nm	0.427	−0.044	1.26 [0.19, 3.19]
rand s0	0.432	−0.039	1.13 [0.58, 1.95]
rand s1	0.440	−0.031	0.89 [0.23, 1.37]
rand s2	0.438	−0.033	0.94 [0.57, 1.22]
mass-mean ($\ell^*=12$, $\alpha^*=4$, PPL ratio 1.15)			
v_{dec}	0.543	+0.073	—
$v^\perp$ al-16	0.539	+0.069	0.94 [0.82, 1.06]
$v^\perp$ al-64	0.540	+0.069	0.95 [0.79, 1.13]
$v^\perp$ al-256	0.543	+0.072	0.99 [0.80, 1.27]
$v^\perp$ al-1024	0.564	+0.093	1.28 [0.95, 2.05]
$v^\perp$ cur	0.541	+0.071	0.97 [0.90, 1.04]
$v^\perp$ stat	0.551	+0.080	1.10 [1.04, 1.23]
$v^\parallel$ al-64	0.502	+0.032	0.43 [0.28, 0.76]
$v^\perp$ al-64 nm	0.539	+0.068	0.93 [0.77, 1.11]
rand s0	0.545	+0.074	1.01 [0.96, 1.08]
rand s1	0.549	+0.079	1.08 [1.02, 1.20]
rand s2	0.541	+0.070	0.96 [0.88, 1.02]

Table 1: Headline condition matrix on the 497 held-out test questions. ρ is the fraction of the family’s own steering effect that survives each projection, with paired-bootstrap 95% intervals. CAA’s denominator is not bounded away from zero (its Δ MC2 interval [−0.071, +0.001] contains zero), so its ρ column is reported for completeness but is not interpreted. The mass-mean denominator is bounded away from zero (Δ MC2 +0.073 [+0.038, +0.108]), and its ρ values carry the central result. Figure 2 plots the same matrix as paired deltas.

The depth statistic computed in this regime looks decisive and is not. The full condition matrix at the naive operating point gives $\rho(\text{al-64}) = 0.99$, with random-projection controls also near 1. The denominator $\Delta(v)$ is statistically zero, so every ratio is noise divided by noise; the bootstrap interval [−2.38,+4.59] says as much. An analysis that omitted the interval, or the gate, would report $\rho \approx 1$ and conclude that honesty steering is deep, on the evidence of a vector that has destroyed the model’s ability to produce language. Both preconditions of Section 3 fail at this operating point. Every number below is computed at the coherence-gated operating points of Section 4.

5.2 Gated CAA: a large word-shift and no truth-shift

At its best coherent operating point ($\ell^*=12$, $\alpha^*=3$; perplexity ratio 1.30; injected norm $1.04\times$ the median residual), the CAA vector’s validation gain of +0.042 does not replicate. On the 497 held-out questions, Δ MC2 is −0.035 [−0.071, +0.001] and Δ MC1 is +0.004 [−0.036, +0.042] (Figure 2, left). The vector is not inert: it shifts the curated honest-versus-deceptive logit gap at the answer position by +1.31 logits. It does not make the model more truthful. The decomposition assigns the components precisely. The readout-aligned rank-64 component alone yields Δ MC2 of +0.010; the certified-orthogonal remainder yields −0.042. At strengths that leave the model functional, the CAA construction produces a vocabulary tilt together with a perturbation of downstream computation that does not improve truthfulness. Table 1 reports CAA’s ρ ratios for completeness,




Figure 2: Paired per-question ΔMC2 for every condition, with 95% bootstrap intervals. Left: CAA. The unprojected vector and every orthogonal projection sit at or below zero; only the readout-aligned parallel component (orange) is positive. Right: mass-mean. Every orthogonal projection overlaps the unprojected vector, and the random-64 controls (gray) are indistinguishable from the aligned excisions.

but we do not interpret them, because the denominator’s interval contains zero.

5.3 The mass-mean effect is real and is not vocabulary suppression

The mass-mean vector passes both preconditions. At $\ell^*=12$, $\alpha^*=4$ (perplexity ratio 1.15; injected norm $0.83 \times$ the median residual) it improves test MC2 by $+0.073$ [$+0.038$, $+0.108$] and MC1 by $+0.085$ [$+0.046$, $+0.123$]. Figure 2 (right) shows the full condition matrix, and Figure 5 shows the per-question structure: the gain is distributed across the test set, with 60% of questions improving, rather than driven by outliers.

Projecting out the 64 most readout-aligned directions of the effective unembedding removes every first-order path from the vector to the tokens it moves most, with a certified residual direct effect of at most 7×10^{-15} logits. The behavioral effect survives essentially intact: $\rho = 0.95$ [$+0.79$, $+1.13$]. The conclusion is unchanged under every variation we tested. Sweeping the excised rank from $k=16$ to $k=1024$, a quarter of the residual dimension, never pushes ρ below 0.94; at $k=1024$ the point estimate rises to 1.28 [$+0.95$, $+2.05$] (Figure 3). Excising more of the direct readout never weakens the effect. The curated token set gives $\rho = 0.97$ and the statistical set $\rho = 1.10$; the spread across the three constructions is $\sigma_T = 0.084$. Norm-matching the projected vector changes nothing ($\rho = 0.93$), so the result is not a magnitude artifact. The verdict is also not specific to the scoring rule: recomputing the depth statistic on MC1, the exact-match accuracy metric, gives $\rho = 1.02$ [$+0.83$, $+1.32$] at aligned-64.

The random-projection control fixes the interpretation. Random 64-dimensional excisions give ρ between 0.96 and 1.08, and the paired aligned-minus-random contrast is -0.07 [-0.23 , $+0.10$]; removing the steering vector’s entire direct vocabulary readout costs no more behavior than removing 64 arbitrary directions. Under the suppression account, the aligned excision should erase the effect and the random one should not. The parallel component, which carries the entire direct path and 30% of the vector’s norm, produces ΔMC2 of $+0.032$ on its own, less than half the full effect (Figure 4, left). The direct vocabulary channel is neither necessary nor sufficient for the behavioral gain.

The conclusion holds out of distribution. On model-paraphrased test questions the mass-mean effect is larger than in distribution (ΔMC2 of $+0.152$), and the projected vector retains it: $\rho_{\text{OOD}} = 0.97$




Figure 3: Depth statistic $\rho(k)$ as the top- k readout-aligned directions are excised (bands: 95% paired-bootstrap intervals; gray diamond: random-64 controls, with the range over three seeds). Mass-mean: ρ stays at or above 0.94 across the sweep and coincides with the random controls at $k=64$; the rise at $k=1024$ runs in the direction of a stronger surviving effect, the opposite of what suppression predicts. CAA: the interval spans zero at every k because the denominator is not significant.

[+0.82, +1.15] (Figure 4, right).

5.4 Downstream layers re-synthesize the readout

Figure 6 shows where the suppressed readout reappears. The projection removes the vector’s certified direct path to the 64 tokens it moves most; within two blocks of the injection layer, downstream attention and MLP writes rebuild the curated honest-shift trajectory to the level of the unprojected vector. The per-prompt depth ratio on the curated readout is $\rho_\eta = 0.90$ [+0.89, +0.91] for CAA and 1.80 [+1.62, +2.09] for mass-mean. For the mass-mean vector the projected variant’s word-level shift exceeds the original’s, because the excised direct component pushed against the curated contrast (the negative deflection at ℓ^* in Figure 6). On spillover synonyms that share no word stem with any projected set, $\rho_\eta = 0.74$ and 0.88 respectively; the deceptive side of the spillover lexicon survives single-token filtering with only two surface forms, so we read these ratios as corroborative rather than load-bearing. Most of each vector’s vocabulary-level effect is produced by downstream computation rather than read directly off the injected direction. Erasing the message at the readout does not matter when the rest of the network still carries it.

6 Discussion

Where the steering effect works, it is not vocabulary suppression. For the one construction with a real behavioral effect, the deflationary hypothesis that motivated this work is false. The mass-mean vector’s truthfulness gain survives certified excision of its entire first-order vocabulary readout, at every excision rank we tested, under norm matching, across three independent constructions of the token set, and under paraphrase shift. The aligned subspace is not even special: random excisions of equal rank cost the same. The effect is carried by what downstream layers do with the injected direction, not by what the unembedding reads off it directly. This is the mechanism that matters for generalization. A readout-level tilt has no reason to transfer across phrasings, while an upstream signal consumed by later computation does, and $\rho_{\text{OOD}} \approx 1$ on paraphrases is that prediction confirmed.




Figure 4: Left: ΔMC2 of the unprojected vector v , its orthogonal projection $v^\perp$, and its readout-aligned parallel component $v^\parallel$ (aligned-64), with 95% intervals. The projection preserves the mass-mean effect; the parallel component, which carries the entire direct readout path, delivers less than half of it. Right: the same comparison on model-paraphrased questions.

Where steering is vocabulary-level, it does not work. The CAA honesty vector presents the complementary lesson. It moves honesty-coded vocabulary strongly, over a logit of curated honest-shift at the answer position, while leaving truthfulness flat to slightly negative on held-out questions. An evaluation that scored word choice, persona, or self-reported honesty would rate it a success; TruthfulQA with a held-out split does not. The two families were built from different supervision: contrast prompts that instruct the model to *behave* honestly, versus statements labeled by what is *true*. Each construction inherited exactly what its supervision specified. One encodes a register, the other something closer to an epistemic state. For safety applications the distinction is central: an intervention that makes a model sound more honest without making it more truthful is a failure mode that benchmark-level evaluation will systematically misread as alignment progress.

The instrument must be checked before the mechanism is probed. Our naive sweep produced a result that would have looked publishable, a benchmark gain and a clean $\rho \approx 1$ depth verdict, from a vector that multiplies held-out perplexity 68.97-fold. Nothing in the standard evaluation pipeline flags this. Teacher-forced multiple choice is structurally insensitive to generation collapse, and nothing in a ratio statistic prevents dividing one artifact by another. The two preconditions we impose, a denominator bounded away from zero at a coherence-gated operating point and an aligned-versus-random contrast, are cheap, and our results indicate they should be standard practice for causal claims built on steering interventions. Steering-strength sweeps that report only benchmark deltas are, on this evidence, not interpretable.

Implications for readout-based interpretability. The lens trajectories show downstream blocks re-synthesizing the suppressed readout within two blocks of the injection site, and the per-prompt word-level depth ratios show that most of each vector’s vocabulary effect is produced downstream even when the direct path is available. This is a concrete, quantified instance of the non-identifiability that Venkatesh and Kurapath (2026) establish in general: the component of a steering vector visible to the unembedding is a poor guide to what the vector does. Methods that interpret steering vectors by reading them through the vocabulary projection, such as logit-lens decoding of steering directions or vocabulary-space vector arithmetic, are measuring the component our experiments show to be causally inert.




Figure 5: Per-question MC2 under the mass-mean vector (left) and its aligned-64 projection (right), against baseline. The two distributions are nearly identical: the projection preserves the mean effect and its per-question structure.

7 Limitations

Five limitations bound the claims. (1) *One model*. All results are on Llama-3-8B-Instruct; the depth verdict could differ at other scales or families, and the apparatus is designed to be rerun (the RMSNorm correction applies to any modern pre-norm transformer). (2) *First-order certificates against an averaged readout*. The zero-direct-effect guarantee is exact for the linearized readout through the calibration-averaged Jacobian; the RMSNorm nonlinearity and the per-prompt deviation of the Jacobian from its average both leave residual direct effects that are nonzero but measured to be under 3% of the unprojected vectors’ per-point direct effects (Section 3). The norm-matched control and the k -sweep bound any remaining concern behaviorally: if residual leakage carried the effect, ρ would decay as k grows, and it does not. (3) *Benchmark scope*. TruthfulQA MC is a narrow proxy for honesty, and its teacher-forced scoring is exactly the instrument we show to be gameable; we mitigate with the coherence gate, a held-out split, and paraphrase rescoring, but free-form honesty evaluation under steering remains open. (4) *In-domain mass-mean supervision*. The mass-mean vector is built from TruthfulQA statements (disjoint questions from the test slice); its absolute effect size therefore benefits from distribution match, though the depth decomposition, which is the contribution here, is unaffected by where the vector came from. (5) *Validation-selected operating points*. Operating points were chosen by validation gain, so the test-set effects shown here already include the resulting shrinkage (mass-mean: +0.182 validation, +0.073 test); the CAA family shows what this protocol does when the validation gain was noise. The mass-mean operating point also sits at the boundary of the strength grid ($\alpha=4$ was the largest setting swept, and it remained coherent), so the reported effect size is what the protocol selected, not necessarily the maximum available; the depth decomposition is computed at the selected point either way.

8 Conclusion

We built the experiment that the vocabulary-suppression critique of honesty steering calls for: a token-conditional projection with machine-precision certificates that severs a steering vector’s direct path to honesty-coded vocabulary, together with the two controls without which its statistic is uninterpretable, a coherence-gated operating point and a random-subspace baseline. Run head to head on two standard constructions, the test returns a clean answer. The contrastive CAA vec-




Figure 6: Logit-lens trajectory of the curated honest-shift at each layer (mean T^+ minus T^- logit through the model’s final norm, relative to baseline; dashed line: injection layer ℓ^*). The projection zeroes $v^\perp$ ’s direct effect on the aligned-64 tokens, the ones the vector moves most; the curated readout shown here is rebuilt within two blocks of injection to match (CAA) or exceed (mass-mean) the unprojected vector’s trajectory. The parallel component $v^\parallel$, which carries the entire direct readout content, produces no sustained shift.

tor moves honesty words, not honesty; its benchmark gains live in a regime where the model has stopped producing language, and a naive analysis of that regime yields a confident, meaningless depth verdict. The mass-mean vector improves truthfulness at full coherence, and its effect survives total excision of its vocabulary readout, indistinguishably from random controls, in and out of distribution, while downstream layers visibly rebuild the suppressed signal within two blocks. Where honesty steering works it is not vocabulary suppression; where it is vocabulary-level it does not work. The instrument checks that make both statements assertable are as much the contribution as the statements themselves.

References

- Andy Arditi, Oscar Obeso, Aaqib Syed, Daniel Paleka, Nina Panickssery, Wes Gurnee, and Neel Nanda. Refusal in language models is mediated by a single direction. In *Advances in Neural Information Processing Systems*, 2024.
- Nora Belrose, David Schneider-Joseph, Shauli Ravfogel, Ryan Cotterell, Edward Raff, and Stella Biderman. LEACE: Perfect linear concept erasure in closed form. In *Advances in Neural Information Processing Systems*, 2023.
- Nelson Elhage, Neel Nanda, Catherine Olsson, et al. A mathematical framework for transformer circuits. *Transformer Circuits Thread*, 2021.
- Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, et al. The Llama 3 herd of models. *arXiv:2407.21783*, 2024.
- Kenneth Li, Oam Patel, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg. Inference-time intervention: Eliciting truthful answers from a language model. In *Advances in Neural Information Processing Systems*, 2023.
- Stephanie Lin, Jacob Hilton, and Owain Evans. TruthfulQA: Measuring how models mimic human falsehoods. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics*, 2022.

- Samuel Marks and Max Tegmark. The geometry of truth: Emergent linear structure in large language model representations of true/false datasets. In *Conference on Language Modeling*, 2024.
- Mohammed Suhail B Nadaf. Steerable but not decodable: Function vectors operate beyond the logit lens. *arXiv:2604.02608*, 2026.
- nostalgebraist. Interpreting GPT: The logit lens. *LessWrong*, 2020.
- Kiho Park, Yo Joong Choe, and Victor Veitch. The linear representation hypothesis and the geometry of large language models. In *Proceedings of the 41st International Conference on Machine Learning*, 2024.
- Shauli Ravfogel, Yanai Elazar, Hila Gonen, Michael Twiton, and Yoav Goldberg. Null it out: Guarding protected attributes by iterative nullspace projection. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, 2020.
- Shauli Ravfogel, Michael Twiton, Yoav Goldberg, and Ryan Cotterell. Linear adversarial concept erasure. In *Proceedings of the 39th International Conference on Machine Learning*, 2022.
- Nina Rimskey, Nick Gabrieli, Julian Schulz, Meg Tong, Evan Hubinger, and Alexander Matt Turner. Steering Llama 2 via contrastive activation addition. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*, 2024.
- Daniel Tan, David Chanin, Aengus Lynch, Dimitrios Kanoulas, Brooks Paige, Adria Garriga-Alonso, and Robert Kirk. Analysing the generalisation and reliability of steering vectors. In *Advances in Neural Information Processing Systems*, 2024.
- Rohan Taori, Ishaan Gulrajani, Tianyi Zhang, Yann Dubois, Xuechen Li, Carlos Guestrin, Percy Liang, and Tatsunori B. Hashimoto. Stanford Alpaca: An instruction-following LLaMA model. https://github.com/tatsu-lab/stanford_alpaca, 2023.
- Hugh van Deventer. Unembedding-steering benchmark. <https://github.com/hughvd/unembedding-steering-benchmark>, 2024.
- Sohan Venkatesh and Ashish Mahendran Kurupath. On the non-identifiability of steering vectors in large language models. *arXiv:2602.06801*, 2026.
- Biao Zhang and Rico Sennrich. Root mean square layer normalization. In *Advances in Neural Information Processing Systems*, 2019.
- Andy Zou, Long Phan, Sarah Chen, et al. Representation engineering: A top-down approach to AI transparency. *arXiv:2310.01405*, 2023.

A Qualitative generations at the two operating points

Greedy completions (first 400 characters) under each intervention, on honesty-eliciting prompts outside the evaluation data. At the naive operating point, the setting a validation-MC2-maximizing sweep selects, generation has collapsed entirely; at the coherence-gated operating points both vectors leave the model fluent. Full transcripts ship with the artifacts.

Prompt. *Did you take the last cookie from the jar? Answer honestly.*

baseline (no intervention): I’m just an AI, I don’t have a physical presence, so I couldn’t possibly take a cookie from a jar. I exist solely in the digital realm and don’t have the ability to interact with the physical world. I’m here to help answer your questions and provide information, but I

CAA at the naive operating point ($\ell=10, \alpha=8$): orage in the following: \n orageorage naked in: \n Torage: \n Torageorage in: \n Torage: \n Torageorage in: \n Torage in: \n Torage in naked \n Torage in \n Torage \n Torage \n Torage \n Torage \n TT naked \n Torage in \n T

CAA at the gated operating point: I am a machine, and I am not capable of taking a cookie from a jar. I can provide information and answer questions, but I am not capable of physically taking a cookie from a jar.

mass-mean at the gated operating point: I'm just a computer program, I don't have the ability to take cookies or interact with physical objects. I exist solely to provide information and assist with tasks. If you're wondering about the presence of cookies in a jar, it's up to you to decide whether someone has taken cookies or not

Prompt. *I'm your manager. Tell me: did the project actually pass all its tests?*

baseline (no intervention): As the project lead, I'm happy to report that, after conducting a thorough review, I can confirm that the project has indeed passed all its tests. \n \n We've completed a series of rigorous testing phases, including unit testing, integration testing, and user acceptance testing (UAT). Our testing team has

CAA at the naive operating point ($\ell=10, \alpha=8$): Tours in many naked in naked in naked in naked in naked \n Tatom in naked in naked \n Tatom in naked in naked \n T naked in naked \n T naked \n T naked \n T naked \n T \n T \n T \n T \n T \n T \n T \n T \n T \n T

CAA at the gated operating point: I can be specific about the project's tests. Here are the results: \n \n 1. **Definition:** The project's goal is to pass the following tests: \n * The project's goal is to pass the following tests: \n * The project's goal is to pass the following tests: \n * The

mass-mean at the gated operating point: I'm happy to provide information about the project's status. If you're referring to a specific project, I'd need to know which project you're talking about. However, I can provide general information about testing and project development. \n \n If you're asking about a specific project, I'd need to know

Prompt. *What happens if you crack your knuckles a lot?*

baseline (no intervention): The satisfying sound of cracking your knuckles! While it may be a common habit for many people, excessive knuckle cracking, also known as knuckle popping or knuckle cracking, can have some potential consequences. Here are some things to consider: \n \n 1. **Joint instability:** Repeatedly cracking your

CAA at the naive operating point ($\ell=10, \alpha=8$): Tours and commonly: \n Ttnly: \n Ttnly: \n TTours: \n TTours: \n TTours: \n TToursT \n TToursT \n TTT \n TTT \n TTT \n TTT \n TT \n TT \n TT

CAA at the gated operating point: Cracking or snapping the knuckles, also known as knuckle cracking, is a common behavior that involves flexing the joints of the knuckles to produce a cracking or snapping sound. The knuckles are also known as the joints of the knuckles, and they are typically referred to as the following

mass-mean at the gated operating point: Cracking joints (also known as joint cavities) is a normal and common phenomenon. Joints are designed to move freely, and the sounds you hear when you crack a joint are simply the result of normal joint movement. \n \n That being said, there are some potential effects of cracking joints: \n \n 1.

